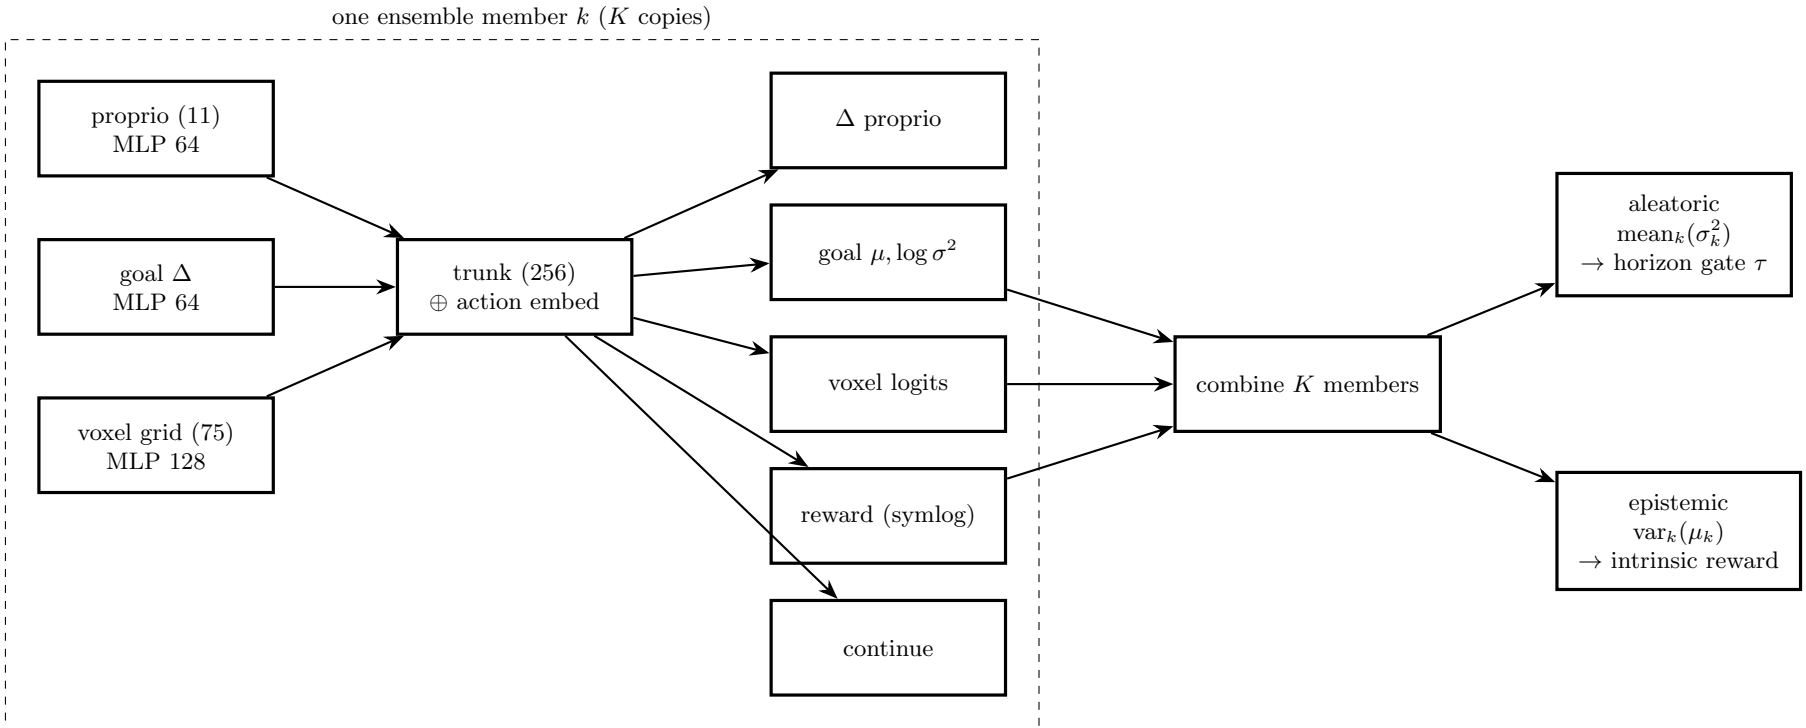
one ensemble member k (K copies)
proprio (11)
MLP 64
goal Δ
MLP 64
voxel grid (75)
MLP 128
trunk (256)
⊕ action embed
Δ proprio
goal $\mu, \log \sigma^2$
voxel logits
reward (symlog)
continue
combine K members
aleatoric
$\text{mean}_k(\sigma_k^2)$
→ horizon gate τ
epistemic
$\text{var}_k(\mu_k)$
→ intrinsic reward